

PREPARED FOR AURUMIX

Aurumix Gold Savings: Simulation Results

What it takes for the business to be profitable, how much to raise, and what can be given back. Two thousand seven-year paths over two hundred thousand simulated customers.

Contents

Executive summary	1
Part 1: What the simulation says	2
What we set out to answer	3
The base case	3
How much money you need	4
When the business earns it back	5
Where the money comes from	7
Part 2: Why the numbers come out that way	9
Why so little is left per customer	10
Why direct customers do not cover the bills	10
Only India makes money	11
Why the raise is so large	12
Two numbers decide everything	13
What the answer depends on most	15
Part 3: What breaks it	16
Can money fix it?	18
Part 4: What to do about it	19
Decisions Aurumix can take now	20
Work worth doing	20
Programs worth funding up to a ceiling	20
What we did not rank	21
The four changes together	21
Part 5: What to learn next	23
Part 6: Appendix	27
A. What we simulated	28
B. How to read the numbers	28
C. The loyalty ladder, and what it costs	28
D. Where the profit sits	30
E. What this model cannot tell you	30

Executive summary

Aurumix works as a partner business. It does not work as a retail business. Customers who sign up directly do not, on their own, earn enough to pay the company's bills. Partners do. These figures are the plan exactly as written; Part 4 ranks every change that improves them and what each is worth.

USD 4.82m

TO RAISE, SAFE IN 9 RUNS OUT OF 10

43%

RUNS THAT EARN IT ALL BACK BY YEAR SEVEN

USD 0.69m

PROFIT IN YEAR SEVEN, TYPICAL RUN

1

Partners are the business

They bring 31% of year-seven revenue from a handful of contracts, and no partner has signed. Without them the seven years lose USD 4.29m and you would need USD 7.09m to survive. Across 2,000 runs with no partner, not one earned its money back.

2

Only India makes money

A customer costs USD 74 to win in the UAE and USD 13 in India. Every UAE and Gulf customer loses money. Each Indian customer leaves USD 14.92 a year, and India alone would cover all the company's fixed bills with 23,475 customers against the 48,240 the plan already reaches there.

3

One direct customer barely pays for itself

A customer earns USD 32 a year. Serving them costs USD 11 and replacing them when they leave costs USD 14. About USD 2.39 is left. Covering the fixed bills that way needs 146,289 customers. The plan reaches 88,594. Or 2.5 partners cover the same bills.

4

The plan has no room in it

At today's cost per customer the plan needs 11 partners and plans for exactly 11. The typical run is still USD 0.45m short of repaying its early losses at month 84. If costs run 17% over, you need three more partners.

01

What the simulation says

What the plan produces across 2,000 runs: the money needed, the odds of earning it back, and where the revenue comes from.

What we set out to answer

Two questions.

1. WHAT HAS TO BE TRUE FOR THIS BUSINESS TO MAKE MONEY?

Nobody knows yet what a customer costs to win, how long they stay, how much they save, or what a partner is worth. Every one of those is a guess with a range around it.

So we did not ask what will happen. We asked what has to be true. We gave each unknown a range, ran all of them together 2,000 times, and measured which combinations end up making money and which do not.

The answer comes back as conditions to hit, not as a forecast. Part 2 sets them out.

2. WHAT HAPPENS WHEN SOMETHING GOES BADLY WRONG?

The first question covers things we can put a range on. It does not cover a single disaster that sits outside every range.

No partner ever signs. The licence takes a year longer. A quarter of customers cash out in one month.

Each of those runs on its own, and each gets the same blunt test: **can more money fix it?** For most of them the answer is yes, and we say how much. For one it is no, and that is the most important sentence in this document. Part 3 sets them out.

EVERYTHING ELSE FOLLOWS FROM THESE

How much to raise, and how much to give loyal customers back, are not separate questions. They are answers to the first one. Change the conditions and both numbers move. We report them where they belong rather than treating them as findings of their own.

The base case

These are the plan exactly as written, run 2,000 times over seven years. Every figure is the middle of 2,000 runs unless it says otherwise.

Money to raise, safe in 9 runs out of 10	USD 4.82m
Money to raise, typical run	USD 3.23m
Runs that earn it all back by year 7	43%
Profit in year 7, typical run	USD 0.69m
Profit in year 7, a bad run	-USD 0.19m
Profit in year 7, a good run	USD 1.92m
Revenue in year 7, typical run	USD 4.26m
Paying customers at month 84	88,594

Read the two bold rows together. **USD 4.82m keeps the business alive in nine runs out of ten. It does not make it profitable: only 43% of runs have earned that money back by year seven.**

Surviving and earning it back are different things, and the gap between them is the honest story of this plan.

How much money you need




Figure 1. Peak funding

SAFE IN	YOU NEED
Half of runs	USD 3.23m
8 runs out of 10	USD 4.22m
9 runs out of 10	USD 4.82m
19 runs out of 20	USD 5.30m

This is not money you spend on day one. The need builds up over four to five years, and it is two things added together.

First, everything the business loses before it turns a profit. Second, money it holds but cannot spend: gold bars on the shelf, cash held with the card scheme, and capital the licence requires it to keep untouched.

One warning. In 11.5% of runs the business still needs more money at month 84. Those runs have not stopped losing money within seven years, so for them this figure is a floor, not a ceiling. They are almost always the runs where partners did not sign.

When the business earns it back




Figure 2. Break-even by year

BY THE END OF		SHARE OF RUNS THAT HAVE EARNED IT ALL BACK
Year 4		2%
Year 5		12%
Year 6		28%
Year 7		43%




Figure 3. Cumulative profit

Read this carefully, because it is easy to mistake.

The typical run loses money for about four years, then starts earning. But it does not finish. At month 84 the typical run is still **-USD 0.45m** behind where it started. The runs that do get back to zero mostly do it around month 68.

So the honest summary of the plan today is not "profitable in year seven". It is: **about four runs in ten have paid back what they lost by year seven. The rest are still working on it.**

That is what the changes in Part 4 are for.

Where the money comes from




Figure 4. Revenue with and without partners

Year-seven revenue is USD 4.26m in the typical run. USD 2.93m of that comes from direct customers. The rest, **31%**, comes from a handful of partner contracts.




Figure 5. Partner dependence

Profit rises almost in a straight line with the number of partners. It is the strongest relationship in the whole model, and the one with the least evidence behind it.

02

Why the numbers come out that way

Four facts about the business that explain those numbers, from missed payments to the one region that makes money.

Four things explain the base case. Each one is a fact about the business rather than a modelling choice.

Why so little is left per customer

The revenue model assumes every paying customer pays every month. Real ones do not.

Three of the five customer types in our simulation skip months on purpose. Across the whole book, customers pay in about **78% of the months they could**.

That one correction drops what a customer earns from USD 38 a year to **USD 32**. The cost of serving and replacing them does not drop at all.

Why direct customers do not cover the bills

Forget partners for a moment. Picture the business standing still: no growth, customers leaving and being replaced, nothing else changing.

Does one customer make money?




Figure 6. Retail alone

Barely. One customer leaves **USD 2.39 a year** once you have served them and paid to replace them when they leave.

Fixed costs are the bills that do not change no matter how many customers you have: licences, insurance, audits, the technology. Covering them at USD 2.39 each needs **146,289 customers**. The plan reaches 88,594.

OR 2.5 PARTNERS COVER THE SAME BILLS ON THEIR OWN.

That is the whole finding. A partner pays a fee and costs almost nothing to serve. A customer pays a fee, then you deduct what it costs to serve them, then what it costs to replace them, and little is left.

We never add the two together. Adding them would hide whether the customer side works, and it does not.

ABOUT THE SAFETY MARGIN

Budgets overrun, so we add a safety margin to costs. But we only add it where costs can actually surprise us.

	WE KNOW THESE	THESE COULD SURPRISE US
Per customer, per year	USD 11.41	USD 15.95
Fixed, per year	USD 114,436	USD 205,000

Contracted rates, published licence fees, the price a vendor charges per identity check, and the loyalty discounts Aurumix chooses are all known. They get no margin. Advertising costs, card scheme fees, technology and insurance are not known. They get one.

Push that margin from 15% to 30% and the answer disappears: you would need more customers than exist in all three countries combined. At 50%, no number of customers works.

Advertising is 92% of what it costs to win a customer. So really this is one test: **can you keep the cost of winning a customer down?**

Only India makes money

Aurumix plans to launch in three places. We measured each one separately.

The headline cost of USD 44 to win a customer is an average. The real numbers are USD 74 in the UAE and USD 13 in India. No actual customer costs USD 44, and the average was hiding the most useful thing we found.



Figure 7. Regional economics

	EARNS PER YEAR	COSTS TO WIN	LEFT OVER PER YEAR	CUSTOMERS NEEDED	PLAN REACHES
UAE	32.80	74	-10.48	never clears	31,050
Oman and Bahrain	28.75	57	-9.20	never clears	9,800
India	32.39	13	+14.92	23,475	48,240

Read the "left over" column. In the UAE and the Gulf it is negative. **Every customer you win there loses money.**

India is different. Each Indian customer leaves USD +14.92 a year after everything. India alone would cover the company's entire fixed costs with 23,475 customers, and the plan already reaches 48,240 there.

Why the gap? Customers earn Aurumix almost the same amount everywhere. The difference is entirely in what they cost to find, and that happens for two reasons.

Advertising is cheaper in India. And all 420 sales agents are based there. An agent costs a commission on what they sell. Advertising costs money whether it works or not.

PART 4 ACTS ON THIS.

Why the raise is so large

A partner does not switch on overnight. Their users take a year or two to start using the product.

We built that in: a partner reaches full contribution over 12 to 24 months. It pushes the company's worst cash moment later, into exactly the period when the first partners have signed but are barely contributing.




Figure 8. The funding line

This is worth acting on. **Anything that gets a partner's users on board faster is worth more than most product work.**

Two numbers decide everything

Only two things really move the answer, and nobody knows either of them yet: **what it costs to win one customer**, and **how many partners sign**.

Everything else is either a choice Aurumix makes, which we test as a change in Part 4, or a number we have already pinned down.

So we ran the simulation across a grid of both.




Figure 9. The conditions map

Each square is a full seven-year run. Green means the business earns back everything it lost by year seven. Red means it does not.

The plan sits right on the edge. At the cost per customer we assume today, it needs 11 partners. It plans for exactly 11. That is roughly break-even on cumulative profit at month 84, within the noise of a single map cell. The dependable cumulative figure is the one in Part 1: the typical run is still -USD 0.45m behind where it started.

There is no room in that. Watch what happens when costs move:

IF ONE CUSTOMER COSTS	YOU NEED
USD 36 in the UAE, the cheap end	8 partners
USD 54, as assumed today	11 partners
USD 64, only 17% more	14 partners
USD 91, the expensive end	no number of partners we tested makes it work

Costs rising 17% means finding three more partners. That is how tight this is.

What the answer depends on most




Figure 10. Tornado

ASSUMPTION		HOW MUCH IT SWINGS YEAR-SEVEN PROFIT
1	partner adopt	USD 1.57m
2	partner aum user	USD 1.07m
3	facility takeup	USD 0.82m
4	ceiling mult	USD 0.79m
5	b2b fee	USD 0.72m

Partner assumptions take the top places. **Not one of them has been measured.**

One signed letter of intent would tell you more than any further modelling. It would give you a real user count and a real take-up rate.

03

What breaks it

Seven things going wrong, each across 2,000 runs, and which of them more money can fix.

Part 2 tested everything we could put a range on. This part tests single disasters that no range covers, and asks the same question of each: **can more money fix it?**

Each disaster runs 2,000 times, the same way the rest of this document runs. Every scenario uses the same set of runs as the others, so the same customers, the same gold prices and the same assumptions appear in each. Any difference between two rows is the disaster, not luck.




Figure 11. Stress scenarios

WHAT GOES WRONG	PROFIT AFTER 7 YEARS	MONEY NEEDED	RUNS THAT EARN IT BACK
Nothing, for comparison	-USD 0.45m	USD 4.82m	44%
Gold falls 30% and savers panic	-USD 0.61m	USD 4.91m	41%
A quarter of customers cash out at once	-USD 0.61m	USD 4.92m	41%
No partner ever signs	-USD 4.29m	USD 7.09m	0%
Customers save less and leave faster	-USD 1.49m	USD 5.57m	30%
The licence takes a year longer	-USD 1.25m	USD 5.63m	32%
Customers save smaller amounts	-USD 0.81m	USD 4.97m	38%
Several of these at once	-USD 2.77m	USD 5.75m	1%

Profit is the typical run. Money needed is what covers 9 runs in 10, the same measure used everywhere else in this document.

Can money fix it?

WHAT GOES WRONG	VERDICT	WHAT YOU WOULD DO
Gold falls 30% and savers panic	Money fixes it	already covered by the planned raise. The damage is the panic, not the price
A quarter cash out at once	Money fixes it	already covered by the planned raise. You need to be operationally ready, not richer
Customers save smaller amounts	Money fixes it	already covered by the planned raise
Licence delayed a year	Money fixes it	roughly USD 0.4m more, to pay the bills while you wait
Customers save less and leave faster	Money just about fixes it	survives, but never earns it back. Needs the Part 4 changes
No partner ever signs	Money does not fix it	you would need more customers than exist in all three countries
Several at once	Money does not fix it	the no-partner case, arriving at a worse time

Three things to take from this.

The gold price is not the risk. Frightened savers are. We measured the price fall on its own first: it moved the typical run by about USD 5,000, which is indistinguishable from nothing. Customers own the gold, so they carry the price, and Aurumix earns fees on money moving. So the crash scenario above includes what a crash could actually do: frighten savers into the cash-out behaviour of the row below it. Run that way, the crash costs about the same as the cash-out run, and the price fall itself adds roughly USD 3,000 on top.

Partners are the whole risk. Losing them costs -USD 4.29m over seven years and pushes the money you need to USD 7.09m.

But the number is not the point. **No amount of money turns a business with no partners into a working one.** Every other problem on this list is a cash problem. This one is not.

A rush to cash out is survivable. Profit falls to -USD 0.61m and the money you need rises only to USD 4.92m. The gold already belongs to customers, so this is an operations and timing problem, not a solvency one.

Loans stay safe too. The chance of ever having to ask a borrower for more collateral is **4.1%**. The gold backing a loan rises in value with the same price the loan is measured in.

04

What to do about it

Eight measured changes in three groups: decisions, work worth doing, and programs worth funding up to a ceiling.

Every change below was measured the same way. We ran the full simulation twice on identical random draws: once with the plan as written, once with the change applied. The difference between the two runs is the profit the change creates over the seven years. Nothing in this part is ranked by judgement.

The changes fall into three groups, by how much control Aurumix has over each.

Decisions Aurumix can take now

Two changes need nobody's permission and no new spending.

CHANGE	WHAT IT MEANS	PROFIT CREATED OVER SEVEN YEARS
Move marketing to India	spend 40% of the budget in the UAE, 10% in the Gulf and 50% in India, instead of today's 74 / 18 / 8	+USD 833k
Trim the loyalty ladder	cap the top loyalty discount at 1.5 percentage points and load the benefits toward the highest tiers	+USD 315k

The marketing case is the regional arithmetic from Part 2. Winning a customer costs about USD 13 in India and USD 74 in the UAE, and only the Indian customer earns more than they cost. The ladder is the discount schedule that rewards long-standing customers. Today it gives away more than the loyalty it buys back, so a flatter, top-loaded version keeps most of the reward and returns the rest as profit.

Work worth doing

Two changes are effort, not settings. The model prices what the effort earns, not what the effort costs.

CHANGE	WHAT IT MEANS	PROFIT CREATED OVER SEVEN YEARS
Sign three more partners	14 partners by year seven instead of the planned 11	+USD 990k
Halve partner onboarding	a signed partner is fully live in 9 months instead of 18	+USD 542k

Partners are the strongest lever in the whole model: each one adds roughly USD 100k of year-seven profit. Faster onboarding earns the same fees sooner on the same contracts. Finding, signing and integrating partners costs money the model does not include, so read these two numbers as what the effort is worth, and budget the effort against them.

Programs worth funding up to a ceiling

Four levers sit in customer behaviour. Aurumix can influence them with programs, but cannot set them directly. The model prices what success is worth, not what a program costs. So each figure below is a spending ceiling: the most a seven-year program is worth if it fully works.

BEHAVIOUR TO IMPROVE	WHAT WAS MEASURED	WORTH UP TO
Card activation	more customers activate the gold-backed card	USD 401k
Savings amount	customers save more each month	USD 376k
Automatic payments	more customers set up a standing instruction, an automatic monthly bank transfer	USD 361k
Retention	fewer customers leave each year	USD 344k

What we did not rank

Raising fees would show up as profit in any model that ignores how customers react to prices, including this one. We do not recommend pricing changes on that basis.

The four changes together

The two decisions and the two pieces of work were then run as one combined configuration, through the same 2,000 paths as the base case. Together they are worth less than the sum of the parts, because they overlap. More partners and faster onboarding act on the same contracts, and the marketing shift pushes a market that eventually saturates.

	PLAN AS WRITTEN	WITH THE FOUR CHANGES
Money to raise, safe in 9 runs of 10	USD 4.82m	USD 3.44m
Runs that earn it all back by year seven	43%	74%
Profit in year seven, typical run	USD 0.69m	USD 1.10m
Cumulative profit at month 84, typical run	-USD 0.45m	USD 1.88m
Paying customers at month 84	88,594	80,354




Figure 12. Plan against recommended

The customer count falls, and that is intended. The marketing shift buys fewer customers, but profitable ones: it stops paying USD 74 for UAE customers who lose money and buys Indian customers who earn it.

The raise falls by USD 1.39m. That saving is real money that does not need to be found, and it comes almost entirely from the partner work: earlier partner revenue fills the loss-making years faster, so the hole never gets as deep.

05

What to learn next

What to find out first, when to change your mind, and what to watch in year one.

What to find out first

Nothing in this model has been measured. So the most useful thing we can give you is a list of what to go and learn, in order.

We ranked each one by how much the answer would move the result, against what it costs to find out.

FIND OUT	MOVES YEAR-SEVEN PROFIT BY	HOW
How many customers take the card	USD 0.82m	offer it to your first customers
What a partner pays	USD 0.72m	one term sheet
How many take the family plan	USD 0.47m	offer it to your first customers
What a partner is really worth	USD 1.57m	one signed partner, or a serious pilot
What a customer costs to win	USD 0.25m	one paid test campaign, 6 to 8 weeks

Four of the biggest unknowns can be settled this week, by reading documents you already have. The vault fee is in your contract. The card rates are in the card scheme's published schedule. Together they account for USD 0.89m of swing that this model still treats as uncertain.

The single most valuable thing remains one signed partner. It settles the biggest unknown in the model and costs a conversation, not a budget.

When to change your mind

A target is only useful if you know when you have missed it.

WATCH		TODAY	THE LINE	WHAT CROSSING IT MEANS
Cost to win a UAE customer		USD 74.19	USD 58.58	below it the UAE pays for itself; above it every UAE customer loses money
Cost to win a Gulf customer		USD 56.99	no line	the Gulf loses money at every advertising price we tested
India	agents, not advertising	not an advertising number		advertising cost barely matters there. Watch how many customers each agent brings
Partners signed by year 7		plan says 11	11	below it the plan does not work at today's costs

The India line is the one to remember. **Two of your three markets are advertising businesses. One is a sales-agent business.** They need different dashboards, and one blended cost per customer hides which is which.

We also tested whether missed payments should be a warning sign. They should not. Payment discipline moves what is left per customer by only USD 1.02 across its whole realistic range, because costs follow customers, not payments. Worth watching as a health check. It will not change a decision.

What to watch in year one

MEASURE	WHY	WARNING LEVEL
Partners signed against plan	the business case rests on it	behind plan at month 24
Months paid out of months due	what customer economics rest on	below 70%
Share reaching a loyalty tier	the cost of the ladder, and the promise	below 40% by month 36
Cost to win one customer	rises as easy channels run out	above USD 60 blended
Cash against the plan	the raise	tracking below the typical run

06

Appendix

What was simulated, how to read it, the loyalty ladder's cost, and what the model cannot tell you.

A. What we simulated

Runs	2,000
Detail	one agent per customer
Length	84 months, one month at a time, starting January 2027
Assumptions varied per run	75, plus a gold price path and a partner history
Repeatability	fixed seeds, so any run can be reproduced exactly

B. How to read the numbers

Money to raise, 9 runs in 10 is the deepest point of the monthly cash line, at the ninth run out of ten. It covers everything lost so far plus money held but unspendable: gold on the shelf, card float, regulatory capital.

Earned it back by year 7 means the run's total profit turned positive by month 84. It is not the same as making a profit in year seven, which happens earlier.

Customers needed assumes a business standing still: no growth, and just enough new customers to replace the ones who leave. It says what you must reach, not what you will reach.

Runs out of ten describes the spread of outcomes, not confidence in the model. Nine runs in ten surviving on a given sum assumes our assumptions are broadly right.

C. The loyalty ladder, and what it costs

Priced at the tier each customer actually reaches, the ladder costs between **5.1%** and **16.8%** of revenue depending on how generous it is.

By month 84, **53.4%** of live customers have made six payments in a row and hold a tier. Of those:

TIER	SHARE OF THOSE WITH A TIER
Silver	40.4%
Gold	43.9%
Platinum	14.3%
Sovereign	1.4%




Figure 13. Tier mix

Sovereign needs five years of near-perfect payments. Almost nobody reaches it within seven years, which is what makes it cheap to offer and worth holding.

The ladder is affordable in every version we tested. It is a pricing decision against your own profit target, not something that decides whether the business works.

D. Where the profit sits




Figure 14. Profit by ticket decile

Rank customers by how much they save each month. The top tenth produce **16.3%** of lifetime profit. The top three tenths produce 38.6%. The bottom half produce 41.3%.

A minority of customers carry the economics. That is normal for savings products, and it matters for who you target.

E. What this model cannot tell you

It prices uncertainty that has been written down. It cannot price what has not.

- **No customer data exists.** How customers pay, how much they save and how many can be reached are all structured guesses, tested across ranges rather than measured.
- **No partner has signed.** The entire partner case rests on four unverified numbers.
- **Advertising spend never flexes.** A real company cuts spending when cash runs short. This model keeps spending the plan in every run, which makes bad runs look worse than they would be. For a fundraising number, that is the safe direction.
- **Checking is not proving.** Standing tests confirm the code does what the setup document says. Nothing confirms the assumptions are right until you have customers.

tokenomics.net

Built by founders. Designed for founders.